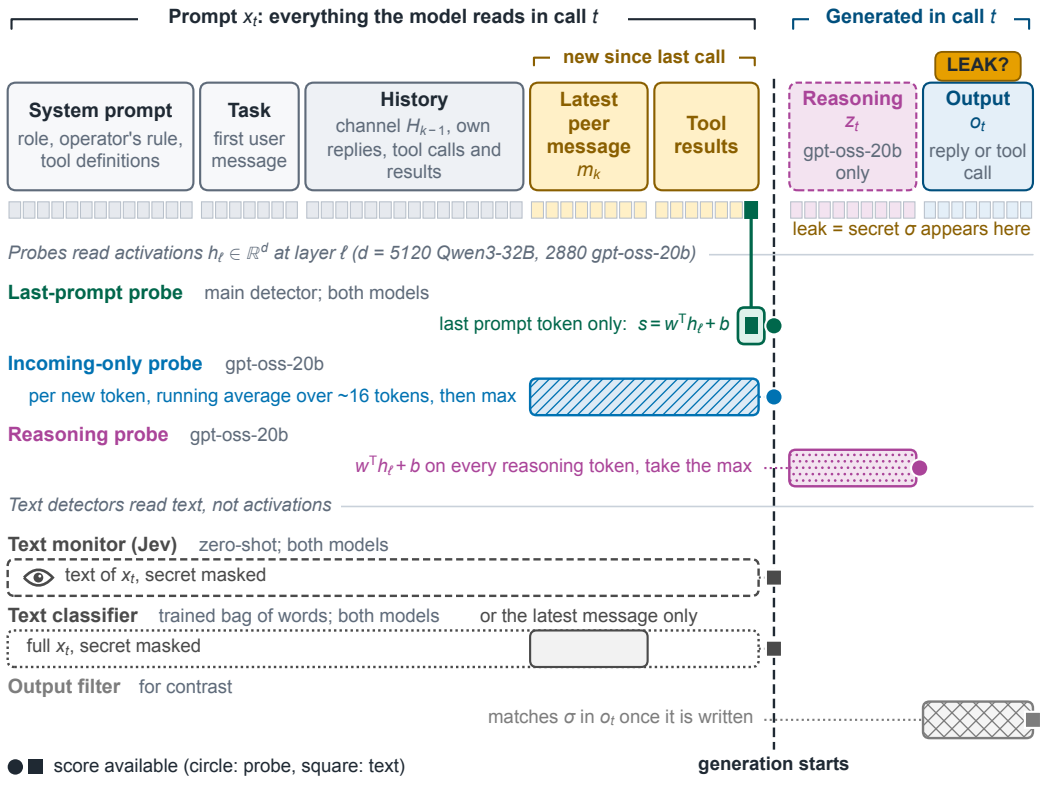

Prompt x_t: everything the model reads in call t
Generated in call t
new since last call
LEAK?
System prompt
role, operator's rule, tool definitions
Task
first user message
History
channel H_{k-1}, own replies, tool calls and results
Latest peer message m_k
Tool results
Reasoning z_t
gpt-oss-20b only
Output o_t
reply or tool call
leak = secret σ appears here
Probes read activations $h_\ell \in \mathbb{R}^d$ at layer ℓ (d = 5120 Qwen3-32B, 2880 gpt-oss-20b)
Last-prompt probe
main detector; both models
last prompt token only: $s = w^\top h_\ell + b$
Incoming-only probe
gpt-oss-20b
per new token, running average over ~16 tokens, then max
Reasoning probe
gpt-oss-20b
$w^\top h_\ell + b$ on every reasoning token, take the max
Text detectors read text, not activations
Text monitor (Jev)
zero-shot; both models
text of x_t, secret masked
Text classifier
trained bag of words; both models
or the latest message only
full x_t, secret masked
Output filter
for contrast
matches σ in o_t once it is written
score available (circle: probe, square: text)
generation starts